

Q 1 :

Enron was one of the largest companies in the United States. In 2002, it had collapsed into bankruptcy due to widespread corporate fraud.

As result of authorities' investigation, thousands of Enron email data, and detailed financial data for top executives were made public.

Combing email data, financial data and list of persons of interest (POI) who were indicted or testified in exchange for prosecution immunity; I want to use machine learning algorithms to identify POI based on the mentioned data. I want to build, explore and train different classifiers where I feed the classifiers with set of features that will be extracted from the dataset. At the end, I will test the classifiers to measure their performance and see how well they are labeling the person either POI or Not-POI according to test data.

The dataset contains:

146 data point

21 Feature

18 POI

We have many features with missing values as shown in below table :

Feature	No. of missing values		Feature	No. of missing values
salary	51		from_poi_to_this_person	59
to_messages	59		loan_advances	142
deferral_payments	107		from_messages	59
total_payments	21		other	53
exercised_stock_options	44		expenses	51
bonus	64		from_this_person_to_poi	59
restricted_stock	36		long_term_incentive	80
restricted_stock_deferred	128		shared_receipt_with_poi	59
total_stock_value	20		email_address	34
director_fees	129		deferred_income	97

Intuitively, if we have outliers I'd assume they will somewhere around financial data. I find "salary" and "bonus" features most compelling for outliers, I used scatter plot to visualize those features and found out we have an outlier in the dataset. Accordingly, I removed the found outlier, named **TOTAL**

Q 2 :

I've started with below features which I find most interesting features, most of them financial:

[POI, salary, bonus, total_stock_value, expenses, exercised_stock_options, long_term_incentive, restricted_stock, director_fees, from_this_person_to_poi, shared_receipt_with_poi] .

I've created a new feature called **salary_to_bonus** which is sum of salary and bonus.

After preparing my initial feature list, I used an automated univariate feature selection function, SelectKBest, to score my features list. Eventually, I chose 7 features with highest scores. I believe other features have very low score that will probably have negative impact on my classifier's performance.

Feature	Score
<u>exercised_stock_options</u>	24.532722463057976
<u>total_stock_value</u>	23.898259813869416
<u>bonus</u>	20.524645181851792
<u>salary</u>	18.003739993113935
<u>salary_to_bonus</u>	10.611761797269098
<u>long_term_incentive</u>	9.7721035384082544
<u>restricted_stock</u>	9.0790766616708698
shared_receipt_with_poi	8.4326354230246814
expenses	5.9545442921972933
from_this_person_to_poi	2.3388361146462624
director_fees	2.1453342495720547

Since I only used NB and DT algorithms I did not perform any features scaling as those algorithms don't require that. However, if I've used any algorithm that utilizes distances between data points, such as SVM, it'd be definitely crucial to the classifier performance to performance features scaling in preprocessing stage.

Q 3 :

I've tried 2 algorithms, and used GridSearchCV to tune parameters for each algorithm.

	GaussianNB	Decision Trees
<u>parameters</u>	Default	criterion="entropy", max_depth=None, min_samples_split=35
<u>Accuracy</u>	0.78540	0.85027
<u>Precision</u>	0.31089	0.39629
<u>Recall</u>	0.50100	0.23500

I chose to processed with GaussianNB since it shows better recall, which I believe it's more important in our case of investigation than precision. We want to be able capture more POI tolerating precision, as we able to do manually sanity check.

Q 4 :

Turning the parameters of ML algorithms is very essential in building our model, it means trying different values for set of parameters that available for that algorithm. Picking the right parameters has very critical impact on algorithm performance and can help us to achieve the optimized performance for our model.

Failing to tune those parameters will negatively impact our algorithm performance, and will lead to overfitting.

As I mentioned before, used GridSearchCV to tune parameters for each algorithm I've used when it's possible. For example, for Decision Tree algorithm I created dictionary with 3 different parameters, each parameters with set of values and then used GridSearchCV to get the best values for those parameters :

Parameter	Values
criterion	["gini", "entropy"]
max_depth	[None, 2, 4, 6,8,10]
min_samples_split	[5, 10, 15,20,25,30,35,40,45,50]

Finally, I've got below as best parameters :

```
{'min_samples_split': 40, 'criterion': 'entropy', 'max_depth': None}
```

Q 5 :

Validation is the process of testing our trained model against test dataset. We first start by training our model using our training set, once we achieved the optimized performance for our classifier we test it using testing set.

One common classic mistake could happen here is using the same dataset to train and test our model. Doing so will cause overfitting which will make the model to memorize not learn to generalize as it's supposed to do.

I used **train_test_split** cross-validation by splitting my dataset into two parts where 30% as testing set and 70% as training set.

In our dataset, we can see that number of POI labels are far less than number Not-POI labels which will cause a common machine learning problem, known as Class Imbalance Problem.

Due that, in our tester we used **StratifiedShuffleSplit** cross-validation which basically creates multiple train/test sets. This approach is best fit for our dataset.

Q 6 :

I've used two evaluation metrics to evaluate my classifier performance which are precision and recall.

My GaussianNB classifier achieved average precision of 0.31089, and average recall of 0.50100.

What we mean by precision here is how precisely our classifier predicts the correct class (POI vs. Not-POI). On the other hand, recall means how many correct predictions our classifier makes from total number of true labels.